



Resources Allocation:

Optimizing Bike Rebalancing in NYC's Citi Bike System with Reinforcement Learning

Lyric Li, Hongxin Song, Lia Wang, Maggie Xu

25.04.29


Motivation

- Urban bike-sharing systems (e.g., CitiBike) offer convenient transportation but often face station imbalances (too full or too empty).
- **Urban Bike-Sharing Challenges**
 - Stations often become too full or too empty.
 - Leads to user frustration and operational inefficiencies.
 - Traditional methods are static and slow to adapt to real-time changes.
- **Our Goal**
 - Model CitiBike station dynamics using real-world trip data.
 - Train a reinforcement learning agent to learn adaptive rebalancing.
 - Optimize bike availability and improve commuter experience.

Data

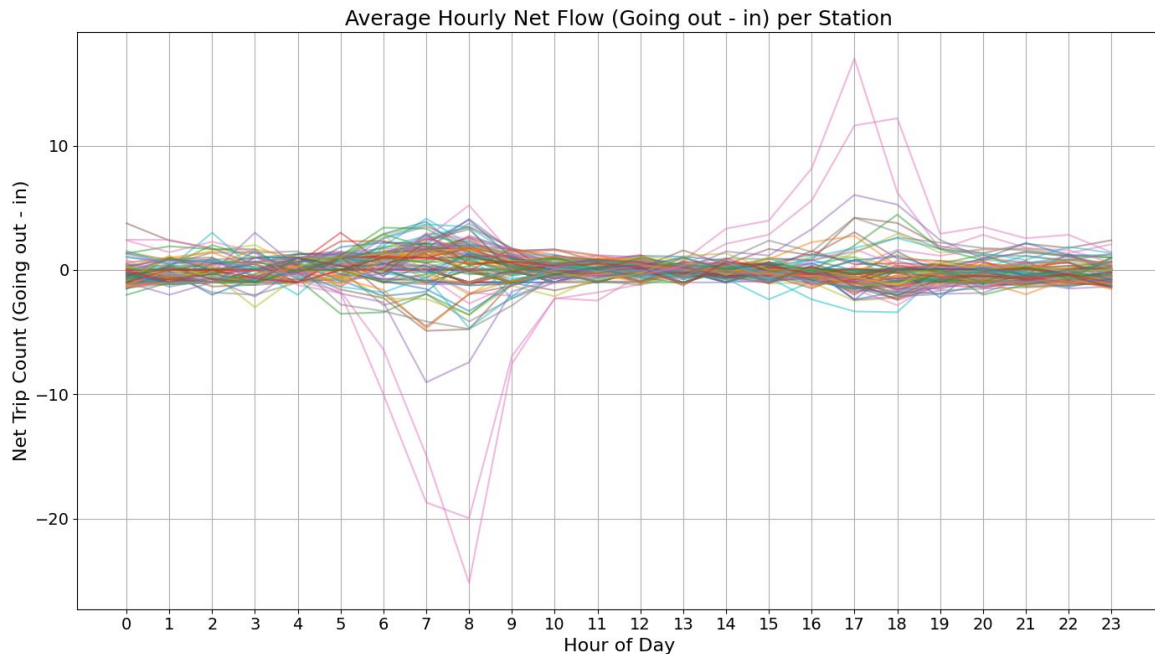
Citi Bike Dataset

- Trip History Data (March 2025):
- Contains **detailed records of ~73,000 rides**, including ride ID, rideable type, start/end times and locations, station names and IDs, geolocation (latitude/longitude), and user type (member or casual).
- Curated to exclude staff, test trips, and rides shorter than 60 seconds.



Clear peaks at 17:00 - 18:00, up to 26 trips in an hour at some stations. Consistent daily outflow of 3 - 5 per hour from midday onward. Environment setted based on that.

Data



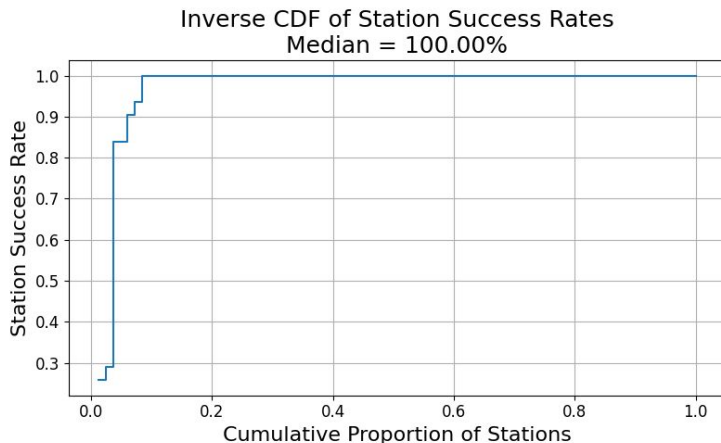
Due to the clear variance between stations, some drain bikes consistently, others accumulate, we segmented stations based on the demand level.

Environment

- Time & Steps:
 - Discrete 2-hour decision intervals
 - 12 steps per 24-hour episode
- Initial Inventory:
 - [40, 30, 20] bikes at $t=0$
- Station Capacity:
 - 60 bikes for high-demand stations
 - 50 bikes for medium-demand stations
 - 40 bikes for low-demand stations
- Natural Refill:
 - Starting from 6:00 AM, +10 bikes are injected every 3 hours to simulate casual returns.
- Rebalancing (Action Space):
 - Every 2 hours
 - Actions: [-5, -3, 0, +3, +5]
 - (Negative = remove bikes, Positive = add bikes)
- State Representation:
 - [bike_count, hour_bucket (0-11), demand_level (one-hot)]
 - Demand levels: High / Medium / Low
- Demand-Level Target Ranges:
 - High: [20, 50] bikes
 - Medium: [10, 40] bikes
 - Low: [5, 30] bikes
- Reward:
 - +10 if stock within target range
 - -10 if outside range
 - -20 if empty/full
 - -0.5 per bike moved

PPO Agent


- **Algorithm:** Proximal Policy Optimization (PPO) via Stable-Baselines3
- **Policy:** MlpPolicy (fully-connected actor-critic network)
- **Actor:** produces a probability distribution over the 5 discrete rebalance actions
- **Critic:** predicts the state-value to guide policy updates
- **Advantage Estimation:** uses Generalized Advantage Estimation for low-variance learning
- **Objective:** clipped surrogate loss to constrain policy updates within a trust region
- **Training:** 240,000 timesteps for a single policy that generalizes across all stations



Success Day: For a given station, if at least 75% of that day's timestamps fall within the target range, the day is counted as a success.

Success Rate: The number of success days for a station / the total number of days (31).


PPO Agent




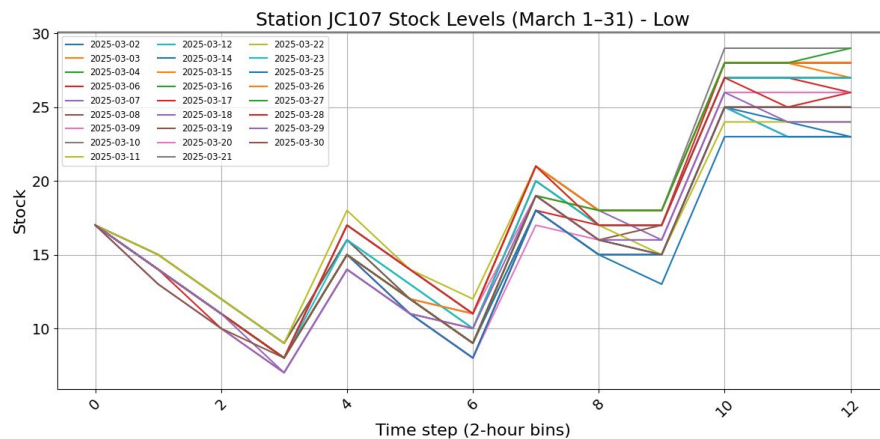
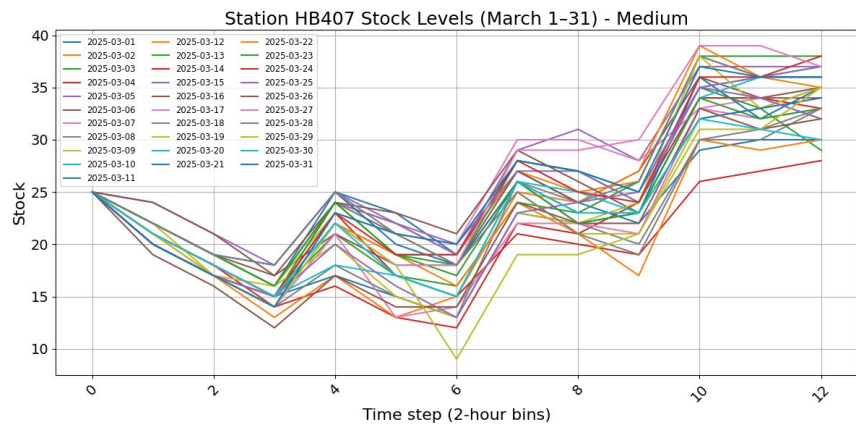
Success Day: For a given station, if at least 75% of that day's timestamps fall within the target range, the day is counted as a success.

Success Rate: The number of stations with a success day / the total number of stations (82).

PPO Agent



PPO Agent



Target Stocks:

- Low: (5, 30)
- Medium: (10, 40)
- High: (20, 50)

Q-Learning

Configuration

- **Q-table:** defaultdict mapping each discretized state to action→value dict
- **Hyperparameters:**
 - Learning rate $\alpha = 0.1$
 - Discount factor $\gamma = 0.95$
 - Episodes = 20 000
- **Exploration:** ϵ -greedy, with ϵ starting at 1.0, decaying by $\times 0.9995$ per episode down to 0.05



Training Loop

- Reset env → observe initial state s
- **While not done:**
 - Pick action a via ϵ -greedy
 - Step, receive reward r and next state s'
 - Update $Q(s,a) \leftarrow Q(s,a) + \alpha [r + \gamma \cdot \max_{a'} Q(s',a') - Q(s,a)]$
 - $s \leftarrow s'$
- **Decay** ϵ after each episode

Policy Evaluation

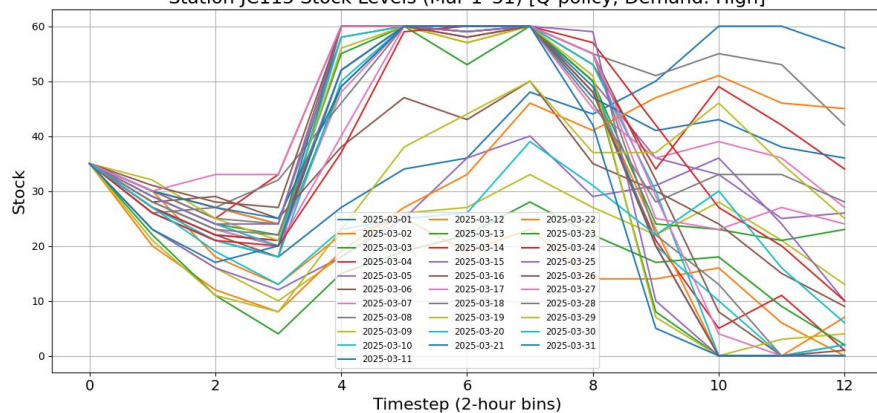
- **For every (station, date) pair:**
 - Instantiate fresh single-station JCEEnvironment
 - Run greedy policy (always pick $\arg \max_a Q(s,a)$) to collect total reward
- Report average episode return

Q-Learning



Q-Learning

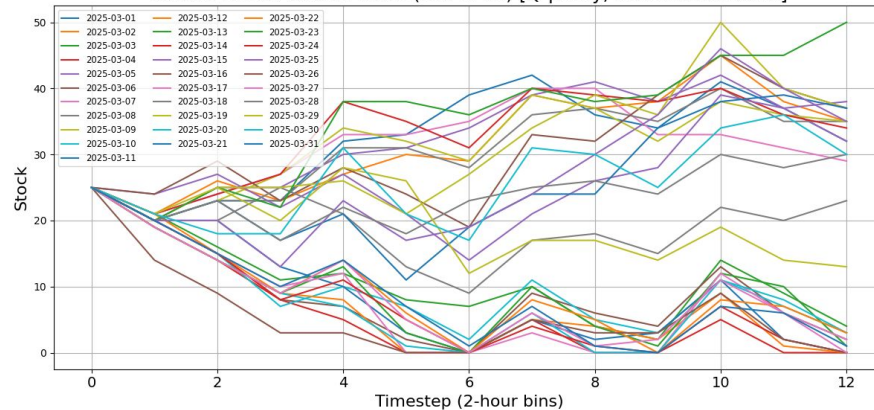
Station JC115 Stock Levels (Mar 1-31) [Q-policy, Demand: High]




Target Stocks:

- Low: (5, 30)
- Medium: (10, 40)
- High: (20, 50)

Station HB407 Stock Levels (Mar 1-31) [Q-policy, Demand: Medium]



Station JC107 Stock Levels (Mar 1-31) [Q-policy, Demand: Low]



Deep Q-Learning with Online Updating

Method Overview:

- Training the DQN on fresh new data day-by-day during March 2025, station by station, while simultaneously updating the network weights
 - “Online”: keep updating after each batch of experiences
- Interact with environment (action → step → observe → reward) in a loop until the day ends.

Action Selection:

- Using an epsilon-greedy policy; Epsilon starts at 0.2 and decays very slightly to encourage gradual exploration → exploitation shift


Experience Collection:

- Every (state, action, reward, next_state, done) tuple is immediately stored into a replay buffer.

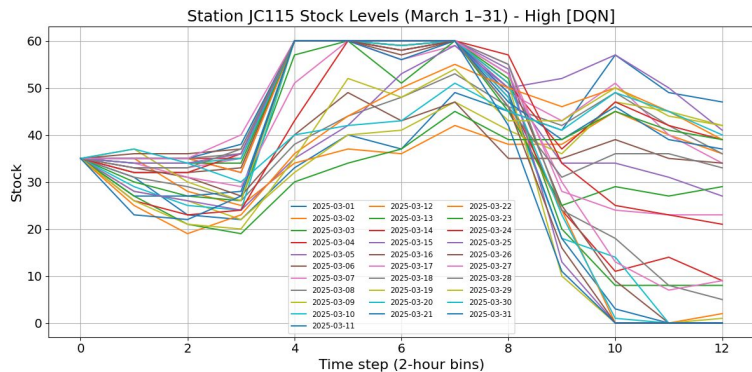
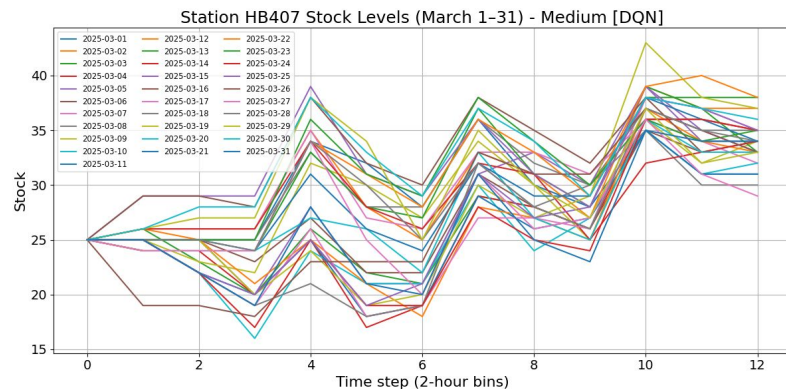
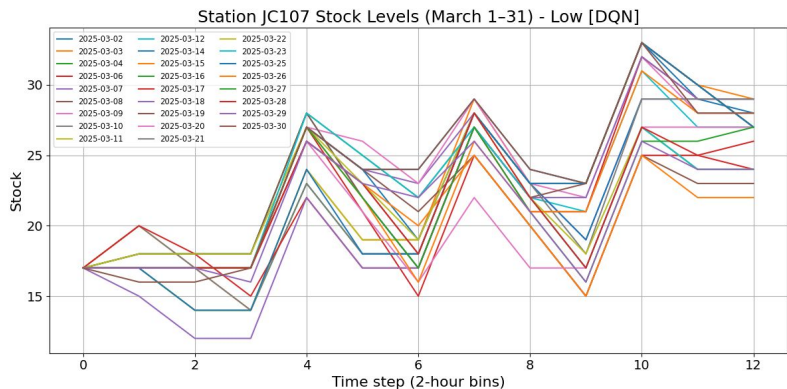
Training & Updating:

- The policy network is updated using sampled experiences with slight epsilon decay after each action, update the policy network weights using Adam optimizer
- Every 50 steps, the target network is synced with the policy network to stabilize training.

Deep Q-Learning with Online Updating




Deep Q-Learning with Online Updating



Target Stocks:

- Low: (5, 30)
- Medium: (10, 40)
- High: (20, 50)


Deep Q-Learning with Online Updating




Conclusion

- Overall, the performances of both PPO and DQN show promising result. PPO achieved the highest station success rate versus Q-Learning and DQN.
- PPO dominates in overall stability without requiring online updates or retraining.
- DQN provides adaptability for dynamic environments but shows more variability (especially for high demand)

PPO



DQN



What Can Be Improved

- **Station-specific Modeling**

- Current approach assigns simple “High/Medium/Low” demand categories.
- Future: model each station’s realistic individual behavior dynamically based on actual flow patterns.

- **Natural Demand Variability**

- Assumes average hourly net flows.
- Future: introduce stochastic demand (simulate daily randomness and peak events like weekends, weather).

- **Policy Generalization**

- Agent is trained and evaluated mostly on March 2025 data.
- Future: test generalization to other months, unseen stations, and different seasonal dynamics.